

Executive Summary

We propose a **standalone AR glasses** platform combining onboard AI vision, speech, and user-personalization to detect traffic violations and support a conversational assistant. The glasses use a high-resolution **microdisplay** (e.g. TI DLP Pico module ¹) with transparent waveguide optics, a suite of sensors (RGB camera, depth/ToF, IMU) and a power-efficient **AR/VR SoC** (e.g. Qualcomm Snapdragon XR2/AR1 series ² ³). For example, the INAIR AR glasses (below) illustrate typical temple-mounted optics and compute with bone-conduction audio **【48†】**. On-device **AI pipelines** run computer vision models (object and license-plate detection) and speech models (ASR, TTS/voice-cloning) entirely locally. Violations (speeding, wrong-way, helmet-off, etc.) are detected by onboard video analysis, then shown to the user with an AR overlay; the user confirms before a “ticket” is issued via secure connectivity. The system emphasizes **privacy and security**: all data (video, speech, user profile) remains on-device, encrypted, and only shared with authorities after user consent. Compliance with Indian regulations (camera and data consent, Traffic/MVA rules) is built in. A phased development plan (18–24 months) is outlined, including hardware prototyping, model development, iterative UX testing, and legal review. The **Bill of Materials (BOM)** and software module breakdown show key components and costs, and the project timeline is charted with milestones.



*Figure: Prototype AR glasses form-factor (e.g. INAIR AR headset). Displays are in the front visor (waveguides), compute and battery in the temple arms, with bone-conduction speakers/sensors ¹ **【48†】**.*

Hardware Components

- **Optics/Display:** Transparent waveguide lenses (DigiLens, Lumus, etc.) with a microdisplay module (e.g. Texas Instruments DLP Pico) to project images into the user’s eyes. TI offers DLP chips for AR:

e.g. the DLP2010 (854×480) up to DLP4710 (1920×1080) ¹ . These engines yield high contrast (>1000:1) at low power ⁴ . Modern AR optics (e.g. pancake lenses or diffractive waveguides) allow thin frames.

- **Processing SoC:** A mobile AR/VR SoC integrates CPU, GPU, DSP, and NPU for ML. Qualcomm's Snapdragon XR platforms (XR2, XR2+ Gen2, AR1 Gen1) are designed for glasses ⁵ ³ . For example, an XR2+ Gen2 reference design drives dual 1080p@90Hz displays with 6-DoF tracking ⁶ and includes a FastConnect 6900 module (Wi-Fi 6E, Bluetooth 5.3, <3 ms latency) ⁵ . We would target a similar or newer chip (e.g. a 2026-era XR2 successor or MediaTek AR chip) with 4–8 GB RAM and neural DSP.
- **Cameras and Sensors:** A front RGB camera (e.g. Sony IMX681 12MP sensor) for scene capture ⁷ , plus a depth-sensing module (e.g. STMicro VL53L5CX ToF sensor: 8×8 zone, 65° FOV, 400 cm range ⁸). These enable object/vehicle detection and 3D mapping (SLAM). Inertial sensors (3-axis gyro+accel, e.g. Bosch BMI270) track head pose. Optionally stereo “world” cameras and eye-tracking IR cameras (as in Magic Leap) can improve AR alignment ⁹ .
- **Audio and Haptics:** MEMS microphone array (e.g. Knowles SiSonic Titan/SPH18R or Falcon/SPH11C3 for high SNR ¹⁰) enables voice commands. Tiny speakers or bone-conduction transducers provide spatial audio. A haptic motor (coin or linear resonant) can give vibration alerts.
- **Battery/PMIC:** A lightweight Li-ion battery (700–1500 mAh per temple arm) supplies ~1–2 hours of active use. TI or Analog Devices power-management ICs handle charging and regulation ¹¹ . (For prototyping, an external battery pack could suffice.) Weight distribution is crucial, so battery often sits in the temples or external belt.
- **Connectivity:** Wi-Fi 6E and Bluetooth 5.x modules (often integrated with the SoC) for syncing, updates, or forwarding tickets. A cellular modem (5G) could be included if direct authority communication is required.

Software Architecture

- **Operating System:** A real-time capable OS (typically Android AOSP or embedded Linux) manages sensors and display. Magic Leap 2 and Meta Quest OS are Android-based; we would start from AOSP 13–15. The OS boots into a locked-down AR environment (“no third-party apps by default”) with a system UI for permissions. All apps run locally; network access is strictly controlled.
- **Firmware and Drivers:** Low-level firmware for cameras, display controller, IMU, etc., provided by SoC vendors. ARCore or OpenXR APIs handle head tracking/SLAM. Drivers interface with the TI DLP projector or similar, and with the microphone/speaker hardware.
- **AR Engine:** A SLAM module fuses camera+IMU to track position and map the scene ¹² . This anchors virtual overlays in 3D (avoiding jitter). We would leverage Google's ARCore or an open SLAM engine (e.g. ORB-SLAM modified for stereo+IMU). SLAM is “perhaps the most important component” of AR software ¹² . On top of SLAM, an object-recognition layer tags real-world objects in view (e.g. vehicles, pedestrians).
- **Vision/ML Stack:** Use an on-device ML runtime (TensorFlow Lite or ONNX Runtime) to run models. A typical pipeline: **Camera → Preprocessor (resize, normalize) → Neural Net → Postprocess**. The OS handles threading, with a real-time priority given to vision inference. GPU/DSP acceleration (via NVIDIA CUDA/TensorRT if PC-connected, or Hexagon/DSP on mobile chips) is used for speed.
- **Speech Stack:** An offline Automatic Speech Recognition (ASR) engine listens for hotwords or queries. Options include Mozilla DeepSpeech/Coqui or a tiny Whisper model. A text-to-speech (TTS) engine (e.g. Tacotron/WaveRNN or Coqui TTS) produces a customizable voice. A dialogue manager coordinates conversations, adapting replies over time (see below).

- **Local Services:** System services record timestamped logs, manage privacy settings, and encrypt data at rest. A secure enclave or ARM TrustZone handles key storage. **Privacy by design:** all image/audio data are stored locally in encrypted form; nothing is uploaded without explicit user consent.

On-Device AI Models

- **Object & License Plate Detection:** A lightweight object detector (e.g. YOLOv4-tiny, YOLOv5s, or MobileNet-SSD) identifies vehicles, traffic lights, lane markings, helmets, phones, etc. ¹³. Pose estimation (e.g. Google's MoveNet or MediaPipe Pose) could detect gestures or fall risk. For license plates, we use a two-stage approach: first a CNN (e.g. a YOLO variant) finds plate regions, then an OCR model (e.g. EasyOCR or Tesseract) decodes characters. The Glasses then query a local/regional database (VAHAN in India) for owner data. Indian traffic cameras already follow this pipeline: "the camera captures the vehicle... AI checks the violation, ANPR reads the plate and fetches owner details...an automated challan is created" ¹⁴. We adapt similar logic but route it via the user for confirmation.
- **Speech Recognition and Generation:** ASR runs continuously (in low-power mode) to catch commands. We would likely use a quantized speech-to-text model (e.g. Whisper Tiny) to convert voice to text. Responses are generated by an on-device language model (e.g. a distillation of GPT or Mistral-7B) or rule-based logic, then fed to TTS. Voice cloning (e.g. SV2TTS architecture) can let users customize the assistant's voice. These models must be pruned/quantized (INT8) for feasibility; typical size might be 10–50 MB each.
- **Personalization/Behavior Models:** The assistant maintains a user profile (preferences, favorite phrases) that evolves. We could use small neural networks or adaptive logic: e.g. learning from user corrections, commands, or tone over time (similar to how phones adapt voice assistants). This might use a simple reinforcement-learning loop or fine-tune embeddings on-device. The adaptation is slow (weeks/months) to ensure stability. All personalization data stays local.

Data Pipelines and Optimization

- **Data Flow:** The camera stream flows into the vision pipeline with minimal buffering. A trigger (e.g. detecting speed or line-crossing) runs the violation model. If a potential violation is flagged, the system extracts the relevant frames and runs the ALPR OCR step. At the same time, audio from the mic can trigger voice command processing. All inference runs on-device: e.g. a plate recognition inferenced in ~100ms on the NPU.
- **Model Quantization/Pruning:** To fit models on-device, we use TensorFlow Lite quantization or ONNX quantization, reducing 32-bit weights to 8-bit. Pruning removes redundant connections. For instance, a quantized YOLOv5s could drop from 50 MB to ~12 MB with minimal accuracy loss. This yields ~4× storage reduction and similar power savings ⁵ ⁸. Latency measurements (using, e.g., Qualcomm's Snapdragon NPU) will guide trade-offs: smaller models = faster, longer battery; larger = more accurate but heavier.
- **Latency/Power Management:** We balance performance and battery by running "heavy" vision analysis only when needed (e.g. GPS/IMU indicates motion near traffic signals). The baseline state uses low-power SLAM and occasional picture capture (~5 FPS) to conserve energy. The NPU handles ML inferences at ~30–60 fps when active. Thermal management (via passive heatsinking in frames) is included.

Security & Privacy

- **Local Data Storage:** All images, video clips, and transcripts are stored on encrypted flash. We use AES-256 encryption tied to device keys. Personal identifiers (names, license numbers) are held only transiently for OCR; no biometric face data is stored.
- **Consent and Permissions:** In accordance with India's new IT security rules, any camera/microphone access shows a persistent indicator and cannot run unnoticed ¹⁵. On first use, the user must explicitly grant camera/mic permission (with rationale). The device periodically reminds the user to review permissions. Users can easily disable the violation-detection feature or clear data at any time.
- **Data Minimization:** By default, no data is sent out. If the user confirms a ticket, only the offending snapshot and metadata (plate, time, location) is transmitted, encrypted, to a designated authority endpoint. We build in suppression of bystander faces/license plates: e.g. faces are blurred unless needed for AR HUD.
- **Resilience & Safety:** The glasses will “fail safe” by shutting off AR overlays if sensors malfunction or battery is critically low, reverting to transparent mode. A tamper-detect banner (per proposed regulations) warns if unauthorized firmware is installed ¹⁵. All critical system modules are code-signed, and update mechanisms require signed packages.

Legal and Regulatory Considerations

Automated enforcement via personal devices is untested in India, so we assume strict limits: **only the police or state systems can issue legal fines**. Private users cannot unilaterally fine others ¹⁶. In practice, the glasses’ “ticket” would be a *report* sent to authorities for review (possibly under new smart-city traffic rules). Any implementation must comply with the Motor Vehicles Act and IT Acts (privacy). Under India’s Digital Personal Data Protection Act 2023, identifiable vehicle data is personal information requiring consent ¹⁷. Hence, the user-facing UI always confirms before “issuing” a fine.

Cameras in public are legal if privacy safeguards are met. New rules disallow background camera use or recording without notice ¹⁵. The AR glasses will display a constant HUD indicator when the camera is active. Traffic police have been barred from using personal phones for evidence (RTO notice) ¹⁸, underscoring that only sanctioned devices generate admissible evidence. Our design treats the glasses as a **trusted personal device**: any photographic evidence it creates carries a digital timestamp and signature for audit (similar to how Chennai’s CCTV e-challans work ¹⁹). All processing is on-device, addressing privacy norms; collected data is shared only for legitimate traffic enforcement.

User Experience (UX) Flows

- **Permission Flow:** On first startup, the glasses prompt the user (via AR overlay and voice) to grant camera/mic permissions. Once granted, an LED or icon remains lit whenever sensors are active. Users can later revoke any permission via voice menu or companion app.
- **Violation Detection Flow:** Continuously, the camera analyzes traffic. If a rule violation is detected (e.g. running a red light), the glasses overlay a bounding box on the offending vehicle and flash a warning icon. The user hears a chime and voice prompt: “*Violation detected: overspeed/helmet. Issue fine?*” The AR display shows the license number and photo frame. The user can respond by voice (“Yes”/“No”) or use a touch-pad on the frame to confirm. On **Yes**, the system packages the evidence

and GPS stamp, and sends it to the authorities (simulating an e-challan). On **No**, the event is logged locally but no report is sent. This flow mimics roadside AI cameras, which “attach the frame as evidence and create an automated challan” when a violation is seen ¹⁴ .



Figure: Typical roadside AI traffic camera setup. In our system, a similar camera (on the glasses) captures images; detected violations trigger an AR popup for user confirmation ¹⁴ [49†] .

- **Conversational Flow:** The user can invoke the assistant (“Hey Auggie” or similar) to ask questions or control settings. The assistant’s voice and style adapt over weeks to the user’s preferences. During any dialogue, permission prompts (e.g. “Allow storing this info?”) appear only for sensitive actions.
- **Safety & Overrides:** A hard “disable” gesture (like tapping the temple) instantly suspends all AI features and returns the device to passthrough (see-through) mode. Battery warnings are shown prominently.

Failure Modes & Safety

The system includes multiple failsafes: if camera data is lost or model confidence is low, no automated actions occur (the glasses simply notify the user without suggesting a fine). Overheating or low battery auto-quiet the AI pipelines to protect the hardware. AR overlays are rendered with very low persistence to avoid motion sickness. Since full blindness is unsafe, digital overlays use a semi-transparent mode and limit brightness to ambient levels. In practice, the glasses resemble high-end “safety glasses” with clear lenses, ensuring the user’s field of view is never obstructed by AR content.

Testing and Validation Plan

We will **bench-test** all components in lab settings: verify camera calibration, model accuracy on recorded video, and end-to-end latency on developer kits (e.g. Magic Leap or Qualcomm dev headsets). We will

collect a dataset of Indian traffic scenes (various lighting, weather) to refine models. Key metrics: object-detection precision/recall, OCR accuracy, ASR word error rate, and total system delay (target <200ms for alerts). **User studies** will measure distraction and acceptance (e.g. perceived usefulness of AR alerts). We will also conduct privacy/security audits: e.g. verifying that no image data leaks from the device. Finally, a small **field pilot** in a controlled environment (with volunteer drivers and observers) will validate the UX and legal compliance before any real-world deployment.

Bill of Materials (estimates)

Component	Supplier / Part #	Est. Unit Cost (USD)	Notes
AR SoC chipset	Qualcomm Snapdragon XR2+ Gen2	\\$150	Includes CPU/GPU/NPU (multicore)
Microdisplay + Optics	TI DLP4710 + waveguide	\\$40	1080p DLP chip (high-brightness) ¹
RGB Camera	Sony IMX681 (12MP)	\\$20	Compact module ⁷
Depth Sensor (ToF)	ST VL53L5CX	\\$10	8×8-zone ToF, 65° FOV ⁸
IMU	Bosch BMI270 6-axis	\\$2	Accelerometer+Gyro
RAM/Storage	SK Hynix LPDDR4 4GB + 64GB eMMC	\\$15	4GB RAM + flash
Microphones (3×)	Knowles Titan SPH18RILM4H (3 units)	\\$6	High-SNR MEMS (digital I2S) ¹⁰
Speakers (bone cond.)	Audeo/Maxx/Adapt	\\$10	Miniature transducers
Battery	Panasonic Li-ion 2000 mAh cell	\\$8	3.7V nominal per cell
Power PMIC	TI buck-boost & charger ICs	\\$5	Low-power regulators
PCB, connectors, etc.	–	\\$20	Includes sensors/cables
Frame & Waveguide	Custom (DigiLens or similar)	\\$50	Holographic optics + frame hardware
Total (approx)		\\$286	Excl. assembly

(Costs are prototype-level estimates. Vendor datasheets and reference designs guided component choices ¹ ⁸ ¹⁰.)

Software Modules and Components

Module/Component	Functionality	Example Technologies
Vision Pipeline	Camera capture, preprocessing	Android Camera2 API, OpenCV
Object Detection	Identify vehicles, people, signs, etc.	TensorFlow Lite (YOLO, SSD)
License Plate OCR	Plate localization & character recognition	EasyOCR/Tesseract
SLAM/Tracking	6-DoF pose and mapping	ARCore/OpenXR SLAM, ORB-SLAM
Speech Pipeline	Microphone input, voice recognition	Vosk, Whisper, Coqui STT
NLP/Dialog Manager	Conversational AI, context management	On-device LLM (Mistral, LLaMA)
TTS / Voice Synth.	Spoken responses, voice cloning	Tacotron 2, WaveRNN, Coqui
User Interface	AR overlay renderer, gesture input	Unity3D/Unreal XR SDK, OpenXR
Connectivity	Networking, sync with authority servers	MQTT/HTTPS, secure TLS
Security Services	Encryption, authentication, audit logging	AES-256, TPM/TEE, SELinux
Power & OS Services	Battery mgmt, permission prompts, OS tasks	Android/Linux kernel, RTOS

(Module breakdown inspired by TI/Qualcomm AR reference designs 1 5 .)

Development Roadmap

A phased prototype timeline is planned as follows:

gantt		
title AR Glasses Development Roadmap		
dateFormat YYYY-MM		
section Phase 1: Concept & Design		
System architecture & requirements	:	a1, 2026-04, 3M
Select components/suppliers	:	a2, after a1, 2M
Bench prototype setup (dev kits)	:	a3, 2026-07, 2M
section Phase 2: Hardware Integration		
PCB design & fabrication	:	b1, 2026-10, 3M
Sensor/display integration	:	b2, after b1, 2M
Initial assembly & power tests	:	b3, 2027-01, 2M
section Phase 3: Software Development		
OS build & driver enablement	:	c1, 2026-08, 4M
SLAM and vision model porting	:	c2, after c1, 4M
Speech assistant implementation	:	c3, 2027-02, 3M

UI/UX flows & personalization	:c4, after c3, 3M
section Phase 4: Testing & Iteration	
Functional testing (lab)	:d1, 2027-06, 3M
User trials and feedback	:d2, after d1, 4M
Performance optimization	:d3, 2027-11, 3M
section Phase 5: Compliance & Finalize	
Regulatory review (safety/privacy)	:e1, 2028-02, 2M
Production-ready refinements	:e2, 2028-04, 3M

In this plan, overlapping workstreams run in parallel (e.g. SLAM development alongside hardware). Each phase ends with a review meeting. Total prototype development is ~24 months.

Each task above includes detailed engineering milestones. For example, “vision model porting” means retraining/customizing object and license-plate detectors on Indian road images. Compliance review will ensure alignment with India’s traffic laws and data privacy rules.

Sources and References

This design leverages published AR and AI research, vendor documentation, and Indian traffic enforcement practices. Key references include hardware datasheets (TI DLP) ¹, industry blogs on AR architecture ² ³, and Indian traffic AI reports ¹⁴ ¹⁶. Legal and privacy points draw on India’s new phone security rules ¹⁵ and critical discussions of AI challans ¹⁷. Vendor datasheets and reference materials (TI, Qualcomm, STMicro, Knowles) informed component selection and integration strategies ¹ ⁸ ¹⁰. All citations above link directly to source material.

¹ ⁴ TI DLP® Pico™ Technology for AR Glasses (Rev. A)

<https://www.ti.com/lit/pdf/dlpa121>

² ¹¹ ¹² AR Glasses Components: The Intricate Anatomy of Augmented Reality – INAIRSPACE

https://inairspace.com/blogs/learn-with-inair/ar-glasses-components-the-intricate-anatomy-of-augmented-reality?srsltid=AfmBOop7r5JiILQMnplx4S-A5HTY0-IL13ezbj3T5eEVaWq5N8ZM8P_8

³ Top Enterprise XR-Related Developments from CES 2025 | CAIL

<https://www.cail.com/top-enterprise-xr-related-developments-from-ces-2025/>

⁵ ⁶ Qualcomm AR Glasses Reference Design Drops the Tether

<https://www.roadtovr.com/qualcomm-xr2-ar-glasses-reference-design/>

⁷ Xiaomi's AI Smart Glasses: Why the IMX681 Choice Validates Our 3mm Camera Revolution - CAMEMAKE

<https://www.camemake.com/ar-vr-smartglasses-imx681-xiaomi/>

⁸ Datasheet - VL53L5CX - Time-of-Flight 8x8 multizone ranging sensor with wide field of view

<https://www.st.com/resource/en/datasheet/vl53l5cx.pdf>

⁹ Sensors | MagicLeap Developer Documentation

<https://developer-docs.magicleap.cloud/docs/device/hardware/sensors/>

10 Knowles Debuts Trio of SiSonic MEMS Microphones for High-Performance TWS and Hearable Designs | audioXpress

<https://audioxpress.com/news/knowles-debuts-trio-of-sisonic-mems-microphones-for-high-performance-tws-and-hearable-designs>

13 14 17 19 How AI and Smart Cameras Are Changing Traffic Challans

<https://blogs.gomechanic.com/ai-traffic-challan/>

15 India's proposed phone security rules that are worrying tech firms | Reuters

<https://www.reuters.com/world/china/indias-proposed-phone-security-rules-that-are-worrying-tech-firms-2026-01-11/>

16 18 Can traffic police use personal devices to click photos or issue fines? : r/pune

https://www.reddit.com/r/pune/comments/1ldj6rw/can_traffic_police_use_personal_devices_to_click/